

TEXT MINING

Lecture 10

SENTIMENT ANALYSIS

KEUNGOU I KIM

awekim@handong.edu



Sentiment Analysis

- Sentiment
 - a view of or attitude toward a situation or event; an opinion [Oxford Languages]
- Sentiment in text
 - In (unstructured) text data, there is also a sentiment or tone.
 - The tone of the text can be captured either by the term or the context.
 - Ex) “This match was tragic but Lorris was extraordinary.”
 - Ex) “Dier was shocking.”
 - Ex) “Conte and Perisic's Dangerous Cohabitation.”

- Also known as opinion mining
 - Use of natural language processing, text analysis, computational linguistics, and biometrics to systematically identify, extract, quantify, and study affective states and subjective information [Wikipedia]
 - Contextual mining of text which identifies and extracts subjective information in source material, and helping a business to understand the social sentiment of their brand, product or service while monitoring online conversations [towarddatascience.com]
- Sentiment analysis allows us to capture “the emotional tone behind a body of text”

- How to label “sentiment” in such cases?
 - Classifying with emoticons whether a tweet is “positive” or “negative” (Read, 2005; Park & Paroubek, 2010)
ex. (positive) ^^, ^-^, (negative) -_- , --;;;
 - Classifying with the number of review points whether the review is “positive “ or “negative” (Pang et al., 2002)
ex. (positive) ★ ★ ★ ★, (negative) ★ ★ ★
 - Classifying with a politician’s behavior whether he is “positive” or “negative” on certain legitimate (Thomas et al., 2006)

- Subjectivity detection
 - Under the assumption that any assumptions or hypotheses are based on the subjective opinion, sentiment analysis can be applied to determine the subjective opinion.
- Opinion classification
 - Sentiment analysis can be used to determine one's opinion
- Targeted sentiment analysis
 - Allows us to distinguish the writer's sentiment on the target object

- Two approaches
 - Lexicon-based approach: determining sentiment using existing knowledge
 - ML-based (supervised) approach: determining sentiment using contextual knowledge

Sentiment Lexicon Resource

dictionary

- AFINN

- Developed by Finn Årup Nielsen
- Sentiment values with a range of -5 to 5.

Finn Årup Nielsen, "A new ANEW: evaluation of a word list for sentiment analysis in microblogs", Proceedings of the ESWC2011 Workshop on 'Making Sense of Microposts': Big things come in small packages. Volume 718 in CEUR Workshop Proceedings: 93-98. 2011 May. Matthew Rowe, Milan Stankovic, Aba-Sah Dadzie, Mariann Hardey (editors)

word	value	word	value	word	value
<chr>	<dbl>	<chr>	<dbl>	<chr>	<dbl>
1 breathtaking	5	1 some kind	0	1 bastard	-5
2 hurrah	5			2 bastards	-5
3 outstanding	5			3 bitch	-5
4 superb	5			4 bitches	-5
5 thrilled	5			5 cock	-5

- (Bing Liu's) BING

- Positive or negative.
- Opinion word expansion and target extraction through double propagation
- One of the well-known opinion lexicon resources

G Qiu, B Liu, J Bu, C Chen. 2011. Computational linguistics, 37(1): 9-27.

word	sentiment	word	sentiment
<chr>	<chr>	<chr>	<chr>
1 2-faces	negative	1 abound	positive
2 abnormal	negative	2 abounds	positive
3 abolish	negative	3 abundance	positive
4 abominable	negative	4 abundant	positive
5 abominably	negative	5 accessible	positive
6 abominate	negative	6 accessible	positive
7 abomination	negative	7 acclaim	positive
8 abort	negative	8 acclaimed	positive
9 aborted	negative	9 acclamation	positive
10 aborts	negative	10 accolade	positive

- EmoLex
 - NRC Word-Emotion Association Lexicon
 - Sentiment words (ex. Trust, fear, negativity , sadness, anger, surprise, positivity, disgust, joy, anticipation)
 - Extracted sentiment information from crowd-sourcing

Mohammad, S. M. & Turney, P. D. 2013. Crowdsourcing a word-emotion association lexicon, Computational Intelligence, 29(3): 436-465.

word	sentiment	word	sentiment	word	sentiment
<chr>	<chr>	<chr>	<chr>	<chr>	<chr>
1 abacus	trust	1 abandon	fear	1 abandon	sadness
2 abbot	trust	2 abandoned	fear	2 abandoned	sadness
3 absolution	trust	3 abandonment	fear	3 abandonment	sadness
4 abundance	trust	4 abduction	fear	4 abduction	sadness
5 academic	trust	5 abhor	fear	5 abortion	sadness
6 accolade	trust	6 abhorrent	fear	6 abortive	sadness
7 accompaniment	trust	7 abominable	fear	7 abscess	sadness
8 accord	trust	8 abomination	fear	8 absence	sadness
9 account	trust	9 abortion	fear	9 absent	sadness
10 accountability	trust	10 absence	fear	10 absentee	sadness

- Laughran-McDonald sentiment lexicon
 - Specialized in the financial field
 - Positive, negative, litigious, uncertain, constraining, superfluous
 - Not recommended to use for text from other fields

word	sentiment			word	sentiment			word	sentiment
<chr>	<chr>			<chr>	<chr>			<chr>	<chr>
1 able	positive			1 abovementioned	litigious			1 aegis	superfluous
2 abundance	positive			2 abrogate	litigious			2 amorphous	superfluous
3 abundant	positive			3 abrogated	litigious			3 anticipatory	superfluous
4 acclaimed	positive			4 abrogates	litigious			4 appertaining	superfluous
5 accomplish	positive	word	sentiment	5 abrogating	litigious	word	sentiment	5 assimilate	superfluous
6 accomplished	positive	<chr>	<chr>	6 abrogation	litigious	<chr>	<chr>	6 assimilating	superfluous
7 accomplishes	positive	1 abandon	negative	7 abrogations	litigious	abeyance	uncertainty	7 assimilation	superfluous
8 accomplishing	positive	2 abandoned	negative	8 absolve	litigious	abeyances	uncertainty	8 bifurcated	superfluous
9 accomplishment	positive	3 abandoning	negative	9 absolved	litigious	almost	uncertainty	9 bifurcation	superfluous
10 accomplishments	positive	4 abandonment	negative	10 absolves	litigious	alteration	uncertainty	10 cessions	superfluous
		5 abandonments	negative			5 alterations	uncertainty		
		6 abandons	negative			6 ambiguities	uncertainty		
		7 abdicated	negative			7 ambiguity	uncertainty		
		8 abdicates	negative			8 ambiguous	uncertainty		
		9 abdicating	negative			9 anomalies	uncertainty		
		10 abdication	negative			10 anomalous	uncertainty		

- `textdata::get_sentiments('name')`

- Imports sentiments list
- `afinn` / `bing` / `nrc`

```
install.packages("tidytext")
> install.packages('textdata')
trying URL 'https://cran.rstudio.com/bin/macosx/contrib/
Content type 'application/x-gzip' length 499072 bytes (4
=====
downloaded 487 KB

The downloaded binary packages are in
  /var/folders/_l/bsz42vx93p185vfk77pkvksr0000gn/T
s   library(tidytext)
> get_sentiments('afinn')
Do you want to download:
  Name: AFINN-111
  URL: http://www2.imm.dtu.dk/pubdb/views/publication_det
  License: Open Database License (ODbL) v1.0
  Size: 78 KB (cleaned 59 KB)
  Download mechanism: https

1: Yes
2: No

Selection: 1
```

```
> get_sentiments('afinn')
# A tibble: 2,477 × 2
  word      value
  <chr>    <dbl>
1 abandon    -2
2 abandoned  -2
3 abandons   -2
4 abducted   -2
5 abduction  -2
6 abductions -2
7 abhor      -3
8 abhorred   -3
9 abhorrent  -3
10 abhors    -3
# ... with 2,467 more rows
```

```
> get_sentiments('bing')
# A tibble: 6,786 × 2
  word      sentiment
  <chr>    <chr>
1 2-faces  negative
2 abnormal negative
3 abolish negative
4 abominable negative
5 abominably negative
6 abominate negative
7 abomination negative
8 abort    negative
9 aborted  negative
10 aborts   negative
# ... with 6,776 more rows
```

```
> get_sentiments('nrc')
# A tibble: 13,872 × 2
  word      sentiment
  <chr>    <chr>
1 abacus    trust
2 abandon   fear
3 abandon   negative
4 abandon   sadness
5 abandoned anger
6 abandoned fear
7 abandoned negative
8 abandoned sadness
9 abandonment anger
10 abandonment fear
# ... with 13,862 more rows
```

Since all these lists are a set of “words,” we need to create a word data set.

- Since sentiment sets are case-sensitive, any sentiment sets can be used.
- If possible, using a specific type of sentiment list that matches the context of the text data sample is recommended

```
> get_sentiments('afinn') %>% filter(word=='abandon')
# A tibble: 1 × 2
  word      value
  <chr>    <dbl>
1 abandon    -2

> get_sentiments('bing') %>% filter(word=='abandon')
# A tibble: 0 × 2
# ... with 2 variables: word <chr>, sentiment <chr>

> get_sentiments('nrc') %>% filter(word=='abandon')
# A tibble: 3 × 2
  word      sentiment
  <chr>    <chr>
1 abandon fear
2 abandon negative
3 abandon sadness
|
```

Lexicon-based Sentiment Analysis Process I

- Lexicon-based sentiment analysis
 - Analyzing the “tone” of the document based on the list of words and their frequency
 - Simply speaking, lexicon-based sentiment analysis is conducting a text frequency analysis with the sentiment lexicon.
 - To do so, a sentiment lexicon resource is needed.

Step 1) Import data

- Analyze Jane Austen's famous novel "Pride & Prejudice"
- Import the data
 - `janeaustenr::austen_books()` returns Jane Austen's six major novels
 - Jane Austen's famous novels are presented in a `data.frame` format → It will make our life much easier!!

```
> library(janeaustenr)
> austen_books() %>% dplyr::select(book) %>% unique
# A tibble: 6 × 1
  book
  <fct>
1 Sense & Sensibility
2 Pride & Prejudice
3 Mansfield Park
4 Emma
5 Northanger Abbey
6 Persuasion
> austen_books() %>% dplyr::select(text) %>% unique %>% data.frame %>%
+   head(30)
```

	text
1	SENSE AND SENSIBILITY
2	
3	by Jane Austen
4	(1811)
5	CHAPTER 1
6	The family of Dashwood had long been settled in Sussex. Their estate
7	was large, and their residence was at Norland Park, in the centre of
8	their property, where, for many generations, they had lived in so
9	respectable a manner as to engage the general good opinion of their
10	surrounding acquaintance. The late owner of this estate was a single
11	man, who lived to a very advanced age, and who for many years of his

Step 2) Text Pre-processing

- Create a variable called *pp* that only contains “Pride & Prejudice”
 - It should not contain any row with missing values
 - Remove the first two rows (Title and author name)

```
> pp <- austen_books() %>%  
+   filter(book == "Pride & Prejudice" & text != "")  
> pp %>% head  
# A tibble: 6 × 2
```

text	book
<chr>	<fct>
1 PRIDE AND PREJUDICE	Pride & Preju...
2 By Jane Austen	Pride & Preju...
3 Chapter 1	Pride & Preju...
4 It is a truth universally acknowledged, that a single man in possession	Pride & Preju...
5 of a good fortune, must be in want of a wife.	Pride & Preju...
6 However little known the feelings or views of such a man may be on his	Pride & Preju...

```
> pp %<>% filter(text != "PRIDE AND PREJUDICE") %>%  
+   filter(text != "By Jane Austen")  
> pp %>% head  
# A tibble: 6 × 2
```

text	book
<chr>	<fct>
1 Chapter 1	Pride & Prej...
2 It is a truth universally acknowledged, that a single man in possession	Pride & Prej...
3 of a good fortune, must be in want of a wife.	Pride & Prej...
4 However little known the feelings or views of such a man may be on his	Pride & Prej...
5 first entering a neighbourhood, this truth is so well fixed in the minds	Pride & Prej...
6 of the surrounding families, that he is considered the rightful property	Pride & Prej...

Step 2) Text Pre-processing

- Create a new column called “ch”
 - Try

```
> pp %>% head
```

```
# A tibble: 6 x 3
```

text	book	ch
<chr>	<fct>	<dbl>
1 It is a truth universally acknowledged, that a single man in possession	Pride & Prejudice	1
2 of a good fortune, must be in want of a wife.	Pride & Prejudice	1
3 However little known the feelings or views of such a man may be on his	Pride & Prejudice	1
4 first entering a neighbourhood, this truth is so well fixed in the minds	Pride & Prejudice	1
5 of the surrounding families, that he is considered the rightful property	Pride & Prejudice	1
6 of some one or other of their daughters.	Pride & Prejudice	1

Step 2) Text Pre-processing

- Create a new column called “ch”
 - “ch” represents the chapter number
 - Check the data and write down codes for generating this new column

```
> pp %>% head
# A tibble: 6 × 2
  text
<chr>
```

*Chapter numbers are included
in the text column*

```
1 Chapter 1
2 It is a truth universally acknowledged, that a single man in possession
3 of a good fortune, must be in want of a wife.
4 However little known the feelings or views of such a man may be on his
5 first entering a neighbourhood, this truth is so well fixed in the minds
6 of the surrounding families, that he is considered the rightful property
```

```
book
<fct>
Pride & Prej...
Pride & Prej...
Pride & Prej...
Pride & Prej...
Pride & Prej...
Pride & Prej...
```

Find the row index of the case including “Chapter”

```
> pp.ch.index <- which(substr(pp$text,1,7)=="Chapter")
```

```
> pp.ch.index
```

```
[1]      1      80     157     303     399     485     694     876    1060    1211    1420    1568    1627    1772
[15]   1871   2017   2315   2429   2875   3039   3182   3357   3508   3654   3823   3958   4151   4266
[29]   4394   4612   4718   4854   4990   5156   5342   5591   5766   5887   5980   6114   6260   6464
[43]   6624   7050   7247   7399   7651   8001   8195   8393   8587   8763   9027   9282   9427   9637
[57]   9888  10035  10253  10473  10612
```

```
> pp$text[pp.ch.index]
```

```
[1] "Chapter 1" "Chapter 2" "Chapter 3"
[7] "Chapter 7" "Chapter 8" "Chapter 9"
[13] "Chapter 13" "Chapter 14" "Chapter 15"
[19] "Chapter 19" "Chapter 20" "Chapter 21"
[25] "Chapter 25" "Chapter 26" "Chapter 27"
[31] "Chapter 31" "Chapter 32" "Chapter 33"
```

Step 2) Text Pre-processing

- Create a new column called “ch”

```
> pp$ch <- 0
> for(i in 1:length(pp.ch.index)){
+   pp$ch[pp.ch.index[i]:(pp.ch.index[(i+1)]-1)] <- i
+ }
Error in pp.ch.index[i]:(pp.ch.index[(i + 1)] - 1) : NA/NaN argument
> pp$ch <-
+   ifelse(pp$ch==0, max(pp$ch)+1, pp$ch)
> pp %<>%
+   filter(substr(text,1,7)!="Chapter")
> pp %>%
+   filter(ch==max(ch))
# A tibble: 107 x 3
```

Error occurs because it does not cover the last chapter, but it's ok.


text	book	ch
<chr>	<fct>	<dbl>
1 Happy for all her maternal feelings was the day on which Mrs. Bennet got	Pride...	61
2 rid of her two most deserving daughters. With what delighted pride	Pride...	61
3 she afterwards visited Mrs. Bingley, and talked of Mrs. Darcy, may	Pride...	61
4 be guessed. I wish I could say, for the sake of her family, that the	Pride...	61
5 accomplishment of her earnest desire in the establishment of so many	Pride...	61
6 of her children produced so happy an effect as to make her a sensible,	Pride...	61
7 amiable, well-informed woman for the rest of her life; though perhaps it	Pride...	61
8 was lucky for her husband, who might not have relished domestic felicity	Pride...	61
9 in so unusual a form, that she still was occasionally nervous and	Pride...	61
10 invariably silly.	Pride...	61

Step 2) Text Pre-processing

- Create chapter – token table called *pp.clean*

- Space tokenization
- Convert to lower cases
- Remove stopwords ('smart' source)
- Remove punctuation
- Lemmatize

```
> '%ni%' <- Negate('%in%')
> library(textstem)
> pp.clean <- 0
> for (i in 1:length(unique(pp$ch))){
+   temp <- pp %>% filter(ch==i)
+   temp.text <- temp$text %>%
+     str_split(" ") %>% unlist %>% tolower
+   temp.text %<>%
+     str_remove_all("[:punct:]") %>%
+     lemmatize_words
+   temp.text <- temp.text[
+     temp.text %ni% stopwords::stopwords('en', source='smart')]
+   pp.clean <- rbind(pp.clean,
+                     data.frame(ch=i, token=temp.text))
+ }
> pp.clean <- pp.clean[2:nrow(pp.clean),]
> pp.clean %<>%
+   filter(token!="") %>%
+   filter(token!="mr") %>%
+   filter(token!="mrs")
> pp.clean %>% head
  ch      token
1  1      truth
2  1 universally
3  1 acknowledge
4  1      single
5  1          man
6  1 possession
```



Remove extra stopwords

Step 2) Text Pre-processing

- As a final step, create a summarized table called *pp.clean.sum*

- *ch – token – n*

- Try

ch	token	n
<i><dbl></i>	<i><chr></i>	<i><int></i>
1	dear	8
1	bennet	7
1	good	6
1	visit	6
1	bingley	5
1	man	5
1	daughter	4
1	girl	4
1	marry	4
1	single	4

Step 2) Text Pre-processing

- As a final step, create a summarized table called *pp.clean.sum*
 - *ch* – token - *n*

```
> pp.clean.sum <-  
+   pp.clean %>% group_by(ch) %>%  
+   count(token) %>%  
+   arrange(ch, desc(n))  
> pp.clean.sum %>% head(10)
```

```
# A tibble: 10 × 3
```

```
# Groups:   ch [1]
```

	ch	token	n
	<dbl>	<chr>	<int>
1	1	dear	8
2	1	bennet	6
3	1	visit	6
4	1	man	5
5	1	bingley	4
6	1	daughter	4
7	1	girl	4
8	1	marry	4
9	1	single	4
10	1	wife	4

Step 2) Text Pre-processing – tidytext::unnest_token()

- tidytext::unnest_tokens()
 - Split a column into tokens, flattening the table into one-token-per-row.
 - Uses *hunspell_parse_tokenizer*: takes a character vector with text (plain, latex, man, html or xml format), parses out the words and returns a list with incorrect words for each line
 - Easy to implement, but not recommended for advanced work

```
> austen_books() %>%
+   filter(book == "Pride & Prejudice" & text != "")
# A tibble: 10,721 x 2
  text
<chr>
1 "PRIDE AND PREJUDICE"
2 "By Jane Austen"
3 "Chapter 1"
4 "It is a truth universally acknowledged, that a single man in possession"
5 "of a good fortune, must be in want of a wife."
6 "However little known the feelings or views of such a man may be on his"
7 "first entering a neighbourhood, this truth is so well fixed in the minds"
8 "of the surrounding families, that he is considered the rightful property"
9 "of some one or other of their daughters."
10 "\"My dear Mr. Bennet,\" said his lady to him one day, \"have you heard that"
# ... with 10,711 more rows
```

```
> austen_books() %>%
+   filter(book == "Pride & Prejudice" & text != "") %>%
+   unnest_tokens(word, text)
# A tibble: 122,204 x 2
  book      word
<fct>    <chr>
1 Pride & Prejudice pride
2 Pride & Prejudice and
3 Pride & Prejudice prejudice
4 Pride & Prejudice by
5 Pride & Prejudice jane
6 Pride & Prejudice austen
7 Pride & Prejudice chapter
8 Pride & Prejudice 1
9 Pride & Prejudice it
10 Pride & Prejudice is
# ... with 122,194 more rows
```

Lexicon-based Sentiment Analysis Process II



Step 3) Sentiment Analysis with 'afinn'

- Meaning of sentiment value of 'afinn'
 - High and positive sentiment value indicates a “positive” term
 - Low and negative sentiment value indicates a “negative” term

```
> pp.word.affin <- austen_books() %>%  
+   filter(book == "Pride & Prejudice") %>%  
+   unnest_tokens(word, text) %>%  
+   inner_join(get_sentiments('afinn'))  
Joining, by = "word"  
> pp.word.affin %>%  
+   left_join(pp.word.affin %>% count(word))  
Joining, by = "word"  
# A tibble: 7,783 × 4  
  book          word    value     n  
  <fct>        <chr>    <dbl> <int>  
1 Pride & Prejudice good         3   200  
2 Pride & Prejudice want         1    44  
3 Pride & Prejudice dear         2   158  
4 Pride & Prejudice no        -1   490  
5 Pride & Prejudice want         1    44  
6 Pride & Prejudice cried       -2    91  
7 Pride & Prejudice want         1    44  
8 Pride & Prejudice no        -1   490  
9 Pride & Prejudice dear         2   158  
10 Pride & Prejudice delighted     3    23  
# ... with 7,773 more rows
```

Step 3) Sentiment Analysis with 'afinn'

- Meaning of sentiment value of 'afinn'
 - High and positive sentiment value indicates a “positive” term
 - Low and negative sentiment value indicates a “negative” term

```
> pp.word.affin %>% group_by(book, word) %>%  
+   dplyr::summarise(value=sum(value)) %>% arrange(desc(value))  
`summarise()` has grouped output by 'book'. You can override us
```

```
argument.  
# A tibble: 846 × 3
```

```
# Groups:   book [1]
```

	book <fct>	word <chr>	value <dbl>
1	Pride & Prejudice	good	600
2	Pride & Prejudice	great	426
3	Pride & Prejudice	dear	316
4	Pride & Prejudice	love	276
5	Pride & Prejudice	pleasure	276
6	Pride & Prejudice	happy	249
7	Pride & Prejudice	hope	242
8	Pride & Prejudice	happiness	216
9	Pride & Prejudice	better	184
10	Pride & Prejudice	affection	174

```
# ... with 836 more rows
```

```
> pp.word.affin %>% group_by(book, word) %>%  
+   dplyr::summarise(value=sum(value)) %>% arrange(value)  
`summarise()` has grouped output by 'book'. You can overri
```

```
argument.  
# A tibble: 846 × 3
```

```
# Groups:   book [1]
```

	book <fct>	word <chr>	value <dbl>
1	Pride & Prejudice	miss	-566
2	Pride & Prejudice	no	-490
3	Pride & Prejudice	cried	-182
4	Pride & Prejudice	ill	-150
5	Pride & Prejudice	lost	-87
6	Pride & Prejudice	poor	-76
7	Pride & Prejudice	afraid	-74
8	Pride & Prejudice	leave	-62
9	Pride & Prejudice	alone	-60
10	Pride & Prejudice	pain	-56

```
# ... with 836 more rows
```

Step 3) Sentiment Analysis with 'bing'

- 'bing' tells us whether a word is “positive” or “negative”
 - Using this classification of “positive” and “negative,” we can simply distinguish the tone of the word.

```
> pp.word.bing <- austen_books() %>%  
+   filter(book == "Pride & Prejudice") %>%  
+   unnest_tokens(word, text) %>%  
+   inner_join(get_sentiments('bing'))
```

Joining, by = "word"

```
> pp.word.bing %>%  
+   left_join(pp.word.bing %>% count(word))
```

Joining, by = "word"

A tibble: 8,704 × 4

	book <fct>	word <chr>	sentiment <chr>	n <int>
1	Pride & Prejudice	pride	positive	48
2	Pride & Prejudice	prejudice	negative	6
3	Pride & Prejudice	good	positive	200
4	Pride & Prejudice	fortune	positive	39
5	Pride & Prejudice	well	positive	224
6	Pride & Prejudice	rightful	positive	1
7	Pride & Prejudice	impatiently	negative	5
8	Pride & Prejudice	objection	negative	15
9	Pride & Prejudice	enough	positive	106
10	Pride & Prejudice	fortune	positive	39

... with 8,694 more rows

Step 3) Sentiment Analysis with 'bing'

- 'bing' tells us whether a word is "positive" or "negative"
 - Comparing the frequency of positive and negative words
 - Finding the top 5 positive and negative words

```
> pp.word.bing %>% group_by(book, sentiment) %>%  
+   dplyr::summarise(count=length(word))  
`summarise()` has grouped output by 'book'. You can  
override the default grouping with `groups=` argument.
```

```
# A tibble: 2 × 3
```

```
# Groups:   book [1]
```

	book	sentiment	count
	<fct>	<chr>	<int>
1	Pride & Prejudice	negative	3652
2	Pride & Prejudice	positive	5052

```
> pp.word.bing %>% group_by(book, sentiment) %>%  
+   dplyr::summarise(count=length(unique(word)))  
`summarise()` has grouped output by 'book'. You can  
override the default grouping with `groups=` argument.
```

```
# A tibble: 2 × 3
```

```
# Groups:   book [1]
```

	book	sentiment	count
	<fct>	<chr>	<int>
1	Pride & Prejudice	negative	838
2	Pride & Prejudice	positive	592

```
> pp.word.bing %>% group_by(sentiment, word) %>%  
+   dplyr::summarise(count=length(word)) %>% ungroup %>%  
+   group_by(sentiment) %>% arrange(desc(count)) %>% slice(1:5)  
`summarise()` has grouped output by 'sentiment'. You can override  
the default grouping with `groups=` argument.
```

```
# A tibble: 10 × 3
```

```
# Groups:   sentiment [2]
```

	sentiment	word	count
	<chr>	<chr>	<int>
1	negative	miss	283
2	negative	object	48
3	negative	scarcely	45
4	negative	impossible	44
5	negative	poor	38
6	positive	well	224
7	positive	good	200
8	positive	great	142
9	positive	enough	106
10	positive	better	92

Step 3) Sentiment Analysis with 'nrc'

- 'nrc' provides the list of sentiments related to the word.
 - Using the list of sentiments, we can find the most representative sentiment that is being used in the text

```
> pp.word.nrc <- austen_books() %>%  
+   filter(book == "Pride & Prejudice") %>%  
+   unnest_tokens(word, text) %>%  
+   inner_join(get_sentiments('nrc'))
```

Joining, by = "word"

```
> pp.word.nrc %>%  
+   left_join(pp.word.nrc %>% count(word))
```

Joining, by = "word"

A tibble: 29,064 × 4

	book <fct>	word <chr>	sentiment <chr>	n <int>
1	Pride & Prejudice	pride	joy	96
2	Pride & Prejudice	pride	positive	96
3	Pride & Prejudice	prejudice	anger	12
4	Pride & Prejudice	prejudice	negative	12
5	Pride & Prejudice	truth	positive	54
6	Pride & Prejudice	truth	trust	54
7	Pride & Prejudice	possession	anger	36
8	Pride & Prejudice	possession	disgust	36
9	Pride & Prejudice	possession	fear	36
10	Pride & Prejudice	possession	negative	36

... with 29,054 more rows

Step 3) Sentiment Analysis with 'nrc'

- 'nrc' provides the list of sentiments related to the word.

```
> pp.word.nrc %>%  
+ left_join(pp.word.nrc %>% count(word)) %>%  
+ group_by(sentiment) %>% dplyr::summarise(n=sum(n)) %>%  
+ arrange(desc(n))  
Joining, by = "word"  
# A tibble: 10 × 2
```

	sentiment	n
	<chr>	<int>
1	positive	860378
2	trust	704070
3	joy	680636
4	anticipation	668688
5	surprise	404757
6	negative	244577
7	sadness	189333
8	fear	125082
9	anger	85166
10	disgust	72395

```
> pp.word.nrc %>%  
+ left_join(pp.word.nrc %>% count(word)) %>%  
+ group_by(sentiment) %>% dplyr::summarise(n=sum(n)) %>%  
+ arrange(n)  
Joining, by = "word"  
# A tibble: 10 × 2
```

	sentiment	n
	<chr>	<int>
1	disgust	72395
2	anger	85166
3	fear	125082
4	sadness	189333
5	negative	244577
6	surprise	404757
7	anticipation	668688
8	joy	680636
9	trust	704070
10	positive	860378

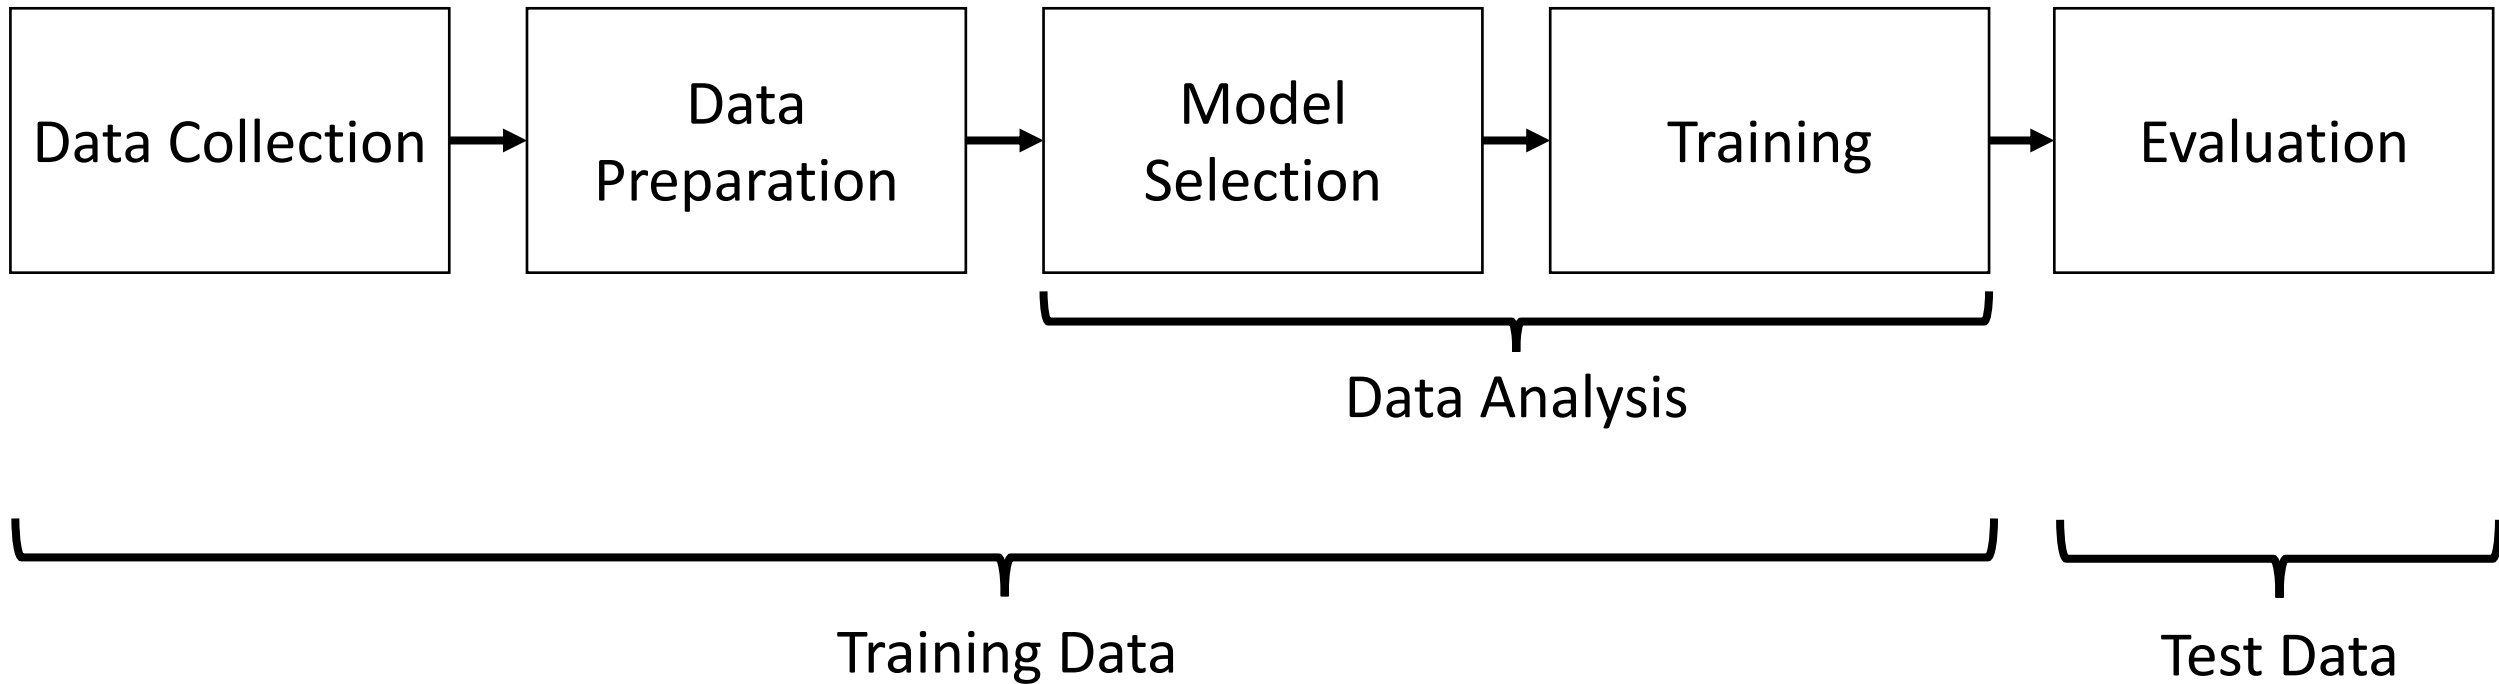
Limitation of Lexicon-based Approach

- In general, a lexicon-based approach is recommended when analyzing a document with a large number of texts.
- If lexicon-based sentiment analysis is used for a short document, it may create an unrealistic result.
 - It's because word frequency matters in a lexicon-based approach
 - An alternative solution is to use an n-gram word.

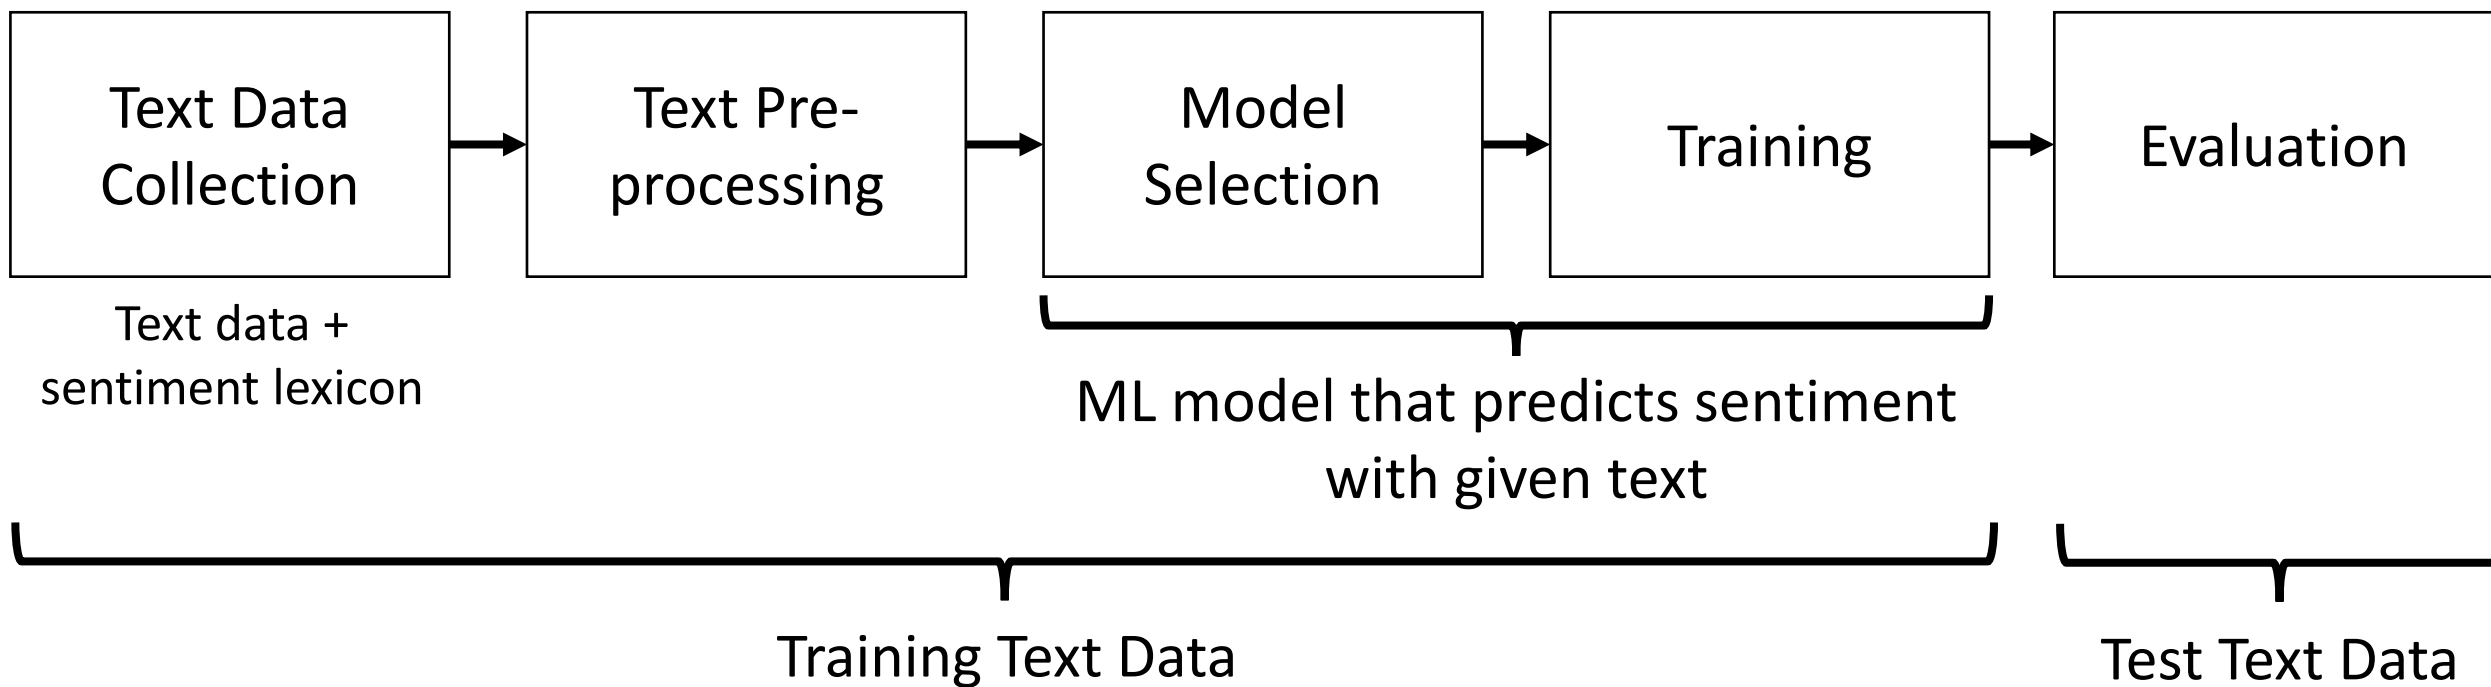
ML-based Sentiment Analysis

- Purpose of machine learning
 - Prediction
 - Classification or regression
- Key concept of machine learning
 - Training data
 - Test data
 - Sampling
 - Label
 - Performance check
 - Validation

- Machine learning process



- ML-based sentiment analysis



- caret (classification and regression training) package
 - R package that allows us to train different types of algorithms using a simple `caret::train()` function
- `caret::train(x, y, method = 'name of ML')`
 - `method = 'rpart' → Decision Tree`
 - `method = 'ranger' → Random Forest`
 - `method = 'xgbTree' → XGBoost`
 - `method = 'knn' → K-Nearest Neighbor`
 - `method = 'nnet' → Neural Network`

ML-based Sentiment Analysis Process

Step 1) Import data

- Create a ML model that predicts the sentiment of text
- Import the data

```
> load(file="R file/R file_LEC10/hs.df.RData")
> hs.df %>% head(1)
```

```
              text
1 I am LOVIN my Life right about now! I'm loving the people God is placing in my life. #Happy&Focu
sed! Striving to be the BEST WOMAN I can be!
```

```
  label type      ID
1 Happy train Happy_1
```

```
> hs.df %>% nrow
```

```
[1] 180
```

```
> hs.df %>% filter(type=="train") %>%
+   filter(label=="Happy") %>% nrow
```

```
[1] 80
```

```
> hs.df %>% filter(type=="train") %>%
+   filter(label=="Sad") %>% nrow
```

```
[1] 80
```

```
> hs.df %>% filter(type=="test") %>%
+   filter(label=="Happy") %>% nrow
```

```
[1] 10
```

```
> hs.df %>% filter(type=="test") %>%
+   filter(label=="Sad") %>% nrow
```

```
[1] 10
```

```
> table(hs.df$label, hs.df$type)
```

	test	train
Happy	10	80
Sad	10	80

Test set

Train set

Step 2) Text Pre-processing

- Create a variable called *hs.df.clean* that only contains cleaned and tokenized text
 - Remember that this is a quick and simple text pre-processing...
- Label – type – ID - word

Remove meaningless numbers, punctuations, etc.

```
> library(dplyr)
> library(stringr)
> library(magrittr)
> library(textstem)
> hs.df.clean <-
+   hs.df %>%
+   unnest_tokens(word, text, "words", to_lower=TRUE) %>%
+   mutate(word = lemmatize_words(word)) %>%
+   anti_join(stop_words, by="word")
> hs.df.clean %>% head
  label type      ID word
1 Happy train Happy_1 lovin
2 Happy train Happy_1 life
3 Happy train Happy_1 love
4 Happy train Happy_1 people
5 Happy train Happy_1 god
6 Happy train Happy_1 life
```

*Each document's
tokenized words listed in
data.frame*

```
> hs.df.clean$word.ed <-
+   hs.df.clean$word %>%
+   str_remove_all("[:digit:]{1,}") %>%
+   str_remove_all("[:lower:]{1,}\\.[[:lower:]]{1,}") %>%
+   str_remove_all("[:punct:]{1,}") %>%
+   str_remove_all("[^a-z0-9]")
> hs.df.clean %<>%
+   filter(word.ed!="")
> dim(hs.df.clean)
[1] 1028    5
```

Step 2) Text Pre-processing

- Create a label vector
 - Train label & Test label
- Create a matrix – train set
 - Document-Term Matrix

```
> hs.df.train.label <-  
+   hs.df$label[hs.df$type=="train"] %>%  
+   as.factor  
> hs.df.test.label <-  
+   hs.df$label[hs.df$type=="test"] %>%  
+   as.factor
```

Any problems?

```
> hs.df.train.dtm <- hs.df.clean %>%  
+   filter(type=="train") %>%  
+   select(ID, word.ed) %>%  
+   count(ID, word.ed) %>%  
+   cast_dtm(document=ID, term=word.ed, value=n)  
> hs.df.train.dtm  
<<DocumentTermMatrix (documents: 160, terms: 481)>>  
Non-/sparse entries: 865/76095  
Sparsity           : 99%  
Maximal term length: 21  
Weighting          : term frequency (tf)  
> hs.df.train.dtmatrix <-  
+   hs.df.train.dtm %>% as.matrix
```

```
> hs.df.test.dtm <- hs.df.clean %>%  
+   dplyr::filter(type=="test") %>%  
+   select(ID, word.ed) %>%  
+   count(ID, word.ed) %>%  
+   cast_dtm(document=ID, term=word.ed, value=n)  
> hs.df.test.dtm  
<<DocumentTermMatrix (documents: 20, terms: 80)>>  
Non-/sparse entries: 110/1490  
Sparsity           : 93%  
Maximal term length: 12  
Weighting          : term frequency (tf)  
> hs.df.test.dtmatrix <-  
+   hs.df.test.dtm %>% as.matrix
```

Step 2) Text Pre-processing

- Create a matrix – test set
 - Document-Term Matrix

```
> hs.df.train.dtmatrix <-  
+   hs.df.train.dtm %>% as.matrix  
> train.vector <- hs.df.clean %>% dplyr::filter(type=="train") %>%  
+   select(word.ed) %>% unique  
> test.vector <- hs.df.clean %>% dplyr::filter(type=="test") %>%  
+   select(word.ed) %>% unique  
> train.test.missing <-  
+   train.vector$word.ed[train.vector$word.ed %ni% test.vector$word.ed]  
> train.test.missing.df <- data.frame(ID=NA,  
+                                     word.ed=train.test.missing)
```

Find list of train-set words that are not included in the test-set

Step 2) Text Pre-processing

- Create a matrix – test set

- Document-Term Matrix

```
> hs.df.test.dtm <- hs.df.clean %>%  
+   dplyr::filter(type=="test") %>%  
+   select(ID, word.ed) %>%  
+   rbind(train.test.missing.df) %>%  
+   count(ID, word.ed) %>%  
+   cast_dtm(document=ID, term=word.ed, value=n)  
> row.names(hs.df.test.dtm)
```

```
[1] "Happy_81" "Happy_82" "Happy_83" "Happy_84" "Happy_85" "Happy_86" "Happy_87" "Happy_88"  
[9] "Happy_89" "Happy_90" "Sad_81" "Sad_82" "Sad_83" "Sad_84" "Sad_85" "Sad_86"  
[17] "Sad_87" "Sad_88" "Sad_89" "Sad_90" NA
```

```
> hs.df.test.dtmatrix <-  
+   hs.df.test.dtm %>% as.matrix  
> hs.df.test.dtmatrix <-  
+   hs.df.test.dtmatrix[1:20,]  
> row.names(hs.df.test.dtmatrix)
```

```
[1] "Happy_81" "Happy_82" "Happy_83" "Happy_84" "Happy_85" "Happy_86" "Happy_87" "Happy_88"  
[9] "Happy_89" "Happy_90" "Sad_81" "Sad_82" "Sad_83" "Sad_84" "Sad_85" "Sad_86"  
[17] "Sad_87" "Sad_88" "Sad_89" "Sad_90"
```

Test set should contain all words that are included in the train set

→ Words that appear in the train set but not in the test set will get value of 0.

Meaningless one added after rbind

Step 3 - RF) ML-based Sentiment Analysis

- Random Forest

```
> library(caret)
> set.seed(1009)
> dim(hs.df.train.dtmatrix)
[1] 160 481
> length(hs.df.train.label)
[1] 160
> dim(hs.df.test.dtmatrix)
[1] 20 517
> length(hs.df.test.label)
[1] 20
> rf.train <- train(x = hs.df.train.dtmatrix,
+                  y = hs.df.train.label,
+                  method = "ranger")
> rf.train
Random Forest

160 samples
481 predictors
 2 classes: 'Happy', 'Sad'

No pre-processing
Resampling: Bootstrapped (25 reps)
Summary of sample sizes: 160, 160, 160, 160, 160, 160, ...
Resampling results across tuning parameters:
```

mtry	splitrule	Accuracy	Kappa
2	gini	0.4703794	0.003731980
2	extratrees	0.4690001	0.001476015
241	gini	0.9748948	0.949440247
241	extratrees	0.9728379	0.945273378
481	gini	0.9740783	0.947757591
481	extratrees	0.9741312	0.947858234

Tuning parameter 'min.node.size' was held constant at a value of 1
Accuracy was used to select the optimal model using the largest value.
The final values used for the model were mtry = 241, splitrule = gini and min.node.size = 1.

*Check the
dimension of data*

*Create a Random
Forest model with
train set*

Step 4 - RF) Performance Check

- Check the performance

- Accuracy
- Kappa
- Sensitivity
- Specificity

```
> rf.train.pred <- predict(rf.train)
> table(rf.train.pred, hs.df.train.label)
      hs.df.train.label
rf.train.pred Happy Sad
      Happy     80   1
      Sad       0  79

> rf.test.pred <-
+   predict(rf.train,
+           newdata = hs.df.test.dtmatrix)
> table(rf.test.pred, hs.df.test.label)
      hs.df.test.label
rf.test.pred Happy Sad
      Happy     10   1
      Sad       0   9
```

*Check the
performance
with the train
set*

*Check the
performance
of model with
the test set*

```
> table(rf.test.pred, hs.df.test.label) %>%
+   confusionMatrix
```

Confusion Matrix and Statistics

	hs.df.test.label	
rf.test.pred	Happy	Sad
Happy	10	1
Sad	0	9

Accuracy : 0.95

95% CI : (0.7513, 0.9987)

No Information Rate : 0.5

P-Value [Acc > NIR] : 2.003e-05

Kappa : 0.9

Mcnemar's Test P-Value : 1

Sensitivity : 1.0000

Specificity : 0.9000

Pos Pred Value : 0.9091

Neg Pred Value : 1.0000

Prevalence : 0.5000

Detection Rate : 0.5000

Detection Prevalence : 0.5500

Balanced Accuracy : 0.9500

'Positive' Class : Happy

Step 3 - NN) ML-based Sentiment Analysis

- Neural Network

```
> nn.train <- train(x = hs.df.train.dtmatrix,  
+                   y = hs.df.train.label,  
+                   method = "nnet")
```

Create a neural
network model

```
# weights: 484  
initial value 110.230278  
iter 10 value 0.005996  
iter 20 value 0.000122  
iter 20 value 0.000047  
iter 20 value 0.000047  
final value 0.000047  
converged
```

```
> nn.train  
Neural Network  
  
160 samples  
481 predictors  
2 classes: 'Happy', 'Sad'
```

```
No pre-processing  
Resampling: Bootstrapped (25 reps)  
Summary of sample sizes: 160, 160, 160, 160, 160, 160, ...  
Resampling results across tuning parameters:
```

size	decay	Accuracy	Kappa
1	0e+00	0.9656416	0.9305773

Step 4 - NN) Performance Check

- Check the performance

- Accuracy
- Kappa
- Sensitivity
- Specificity

```
> nn.train.pred <- predict(nn.train)
> table(nn.train.pred, hs.df.train.label)
      hs.df.train.label
nn.train.pred Happy Sad
      Happy    80   1
      Sad      0  79

> nn.test.pred <-
+   predict(nn.train,
+           newdata = hs.df.test.dtmatrix)
> table(nn.test.pred, hs.df.test.label)
      hs.df.test.label
nn.test.pred Happy Sad
      Happy   10   1
      Sad     0   9
```

*Check the
performance
with the train
set*

*Check the
performance
of model with
the test set*

```
> table(nn.test.pred, hs.df.test.label) %>%
+   confusionMatrix
```

Confusion Matrix and Statistics

	hs.df.test.label	
nn.test.pred	Happy	Sad
Happy	10	1
Sad	0	9

Accuracy : 0.95

95% CI : (0.7513, 0.9987)

No Information Rate : 0.5

P-Value [Acc > NIR] : 2.003e-05

Kappa : 0.9

McNemar's Test P-Value : 1

Sensitivity : 1.0000

Specificity : 0.9000

Pos Pred Value : 0.9091

Neg Pred Value : 1.0000

Prevalence : 0.5000

Detection Rate : 0.5000

Detection Prevalence : 0.5500

Balanced Accuracy : 0.9500

'Positive' Class : Happy